

Data Science for Innovation

W4: Data Mining – Association Rule Mining

Presented by
Ali Anaissi

Overview of Week 4

Today: Data Mining– Association Rule Mining

Objective

Learn techniques for unsupervised learning, with tools in Python.

Lecture

- Association rule mining

Readings

- Intro to Data Mining, Ch. 6
<http://www-users.cs.umn.edu/~kumar/dmbook/ch6.pdf>
- Intro to Data Mining, Ch. 8
<http://www-users.cs.umn.edu/~kumar/dmbook/ch8.pdf>
- Data Science from Scratch, Ch. 11 & 19

Exercises

- Associations from scratch
- Associations using mlxtend
- Associations using pyfpgrowth

Supervised vs. Unsupervised Learning

- Supervised learning (e.g. classification and regression)
 - Supervision: The training data are accompanied by **labels** indicating the class of the observations
- Unsupervised learning (e.g. clustering and association rules)
 - The class labels of training data is **unknown**
 - Given a set of measurements, observations, etc. with the aim of
 - Establishing the existence of classes or clusters in the data
 - Discovering hidden patterns or rules

Unsupervised Learning

- We'll focus on **unsupervised** machine learning techniques
 - *Association rule mining*
 - Dimensionality reduction
 - Outlier detection
 - *Clustering*
 - Etc.

Association Rule Mining

Association Analysis

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

How can businesses
improve sales by
analysing customer
purchase data?

Slides adapted from Tan et al. Introduction to data mining.

[http://www-users.cs.umn.edu/~kumar/dmbook/
http://www-users.cs.umn.edu/~kumar/dmbook/dmslides/chap6_basic_association_analysis.pdf](http://www-users.cs.umn.edu/~kumar/dmbook/http://www-users.cs.umn.edu/~kumar/dmbook/dmslides/chap6_basic_association_analysis.pdf)

Association Rule Mining

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- Predict occurrence of an item based on other items in the transaction, eg:

$\{\text{Diaper}\} \rightarrow \{\text{Beer}\}$

$\{\text{Milk}, \text{Bread}\} \rightarrow \{\text{Eggs}, \text{Coke}\}$

$\{\text{Beer}, \text{Bread}\} \rightarrow \{\text{Milk}\}$

- Note that arrows indicate co-occurrence, not causality

Definition: Itemset

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- An **itemset** is a collection of one or more items
`{Milk,Bread,Diaper}`
- A **k-itemset** is an itemset containing k items

Definition: Frequent Itemset

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- **Support count** (σ) is the itemset frequency
 $\sigma(\{\text{Milk, Diaper, Beer}\}) = 2$
- **Support** (s) is the normalised itemset frequency
$$s = \frac{\sigma(\{\text{Milk, Diaper, Beer}\})}{|T|} = \frac{2}{5}$$
- A **frequent itemset** has $s \geq \text{min_support}$

Definition: Association Rule

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- An **association rule** is an implication of the form $X \rightarrow Y$ where X and Y are itemsets
 $\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$
- **Confidence** (c) measures how often Y occurs in transactions with X
$$c = \frac{\sigma(\{\text{Milk, Diaper, Beer}\})}{\sigma(\{\text{Milk, Diaper}\})} = \frac{2}{3}$$
- An **association rule** has $c \geq \text{min_conf}$

Mining Association Rules

1. Frequent itemset generation

- Generate all itemsets with $s \geq \text{min_support}$

2. Rule generation

- Generate high-confidence rules from each frequent itemset
- Each rule is a binary partitioning of a frequent itemset

Finding Frequent Itemsets

TID	Items
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- Let $min_support = 50\%$
- Freq. 1-itemsets:
 - Beer:3(60%);
Nuts:3(60%);
Diaper:4(80%);
Eggs:3(60%)
- Freq. 2-itemsets:
 - {Beer, Diaper}:3(60%)

Finding Association Rules

TID	Items
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

Market-basket transactions

TID: Transaction Identifier

Items: Transaction item set

- Let $min_support = 50\%$, $min_conf = 50\%$
- Freq. Pat.: Beer:3, Nuts:3, Diaper:4, Eggs:3, {Beer, Diaper}:3
- Association rules:
 - Beer $\rightarrow$ Diaper (100%)
 - Diaper $\rightarrow$ Beer (75%)

Mining Association Rules

1. Frequent itemset generation


- Generate all itemsets with $s \geq \text{min_support}$

2. Rule generation

- Generate high-confidence rules from each frequent itemset
- Each rule is a binary partitioning of a frequent itemset

Easy! But brute force enumerate is computationally prohibitive..

There are 2^d candidate itemsets!



**Enumeration of $2^5 - 1$
candidate itemsets
for {A,B,C,D,E}**

Reducing the number of candidates

- **Apriori Principle**

If an itemset is frequent, then all of its subsets are also frequent


- **Conversely**

If an itemset is infrequent, then its supersets are also infrequent

- The support of an itemset never exceeds the support of its subsets
 - $s(\text{Bread}) \geq s(\text{Bread, Beer})$
 - $s(\text{Milk}) \geq s(\text{Bread, Milk})$
- This is known as the anti-monotone property of support

Pruning the 2^d candidate itemsets

If $s(\{A,B\}) \leq \text{min_support}$



Then we can
prune supersets
from the search
space

Apriori algorithm for generating frequent itemsets

While the list of $(k-1)$ -itemsets is non-empty:

- Generate candidate k -itemsets

- Identify and keep frequent k -itemsets

The Apriori Algorithm—An Example

Min_sup = 2

Database TDB

Tid	Items
1	A, C, D
2	B, C, E
3	A, B, C, E
4	B, E

1st scan

C_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

L_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3

L_2

Itemset	sup
{A, C}	2
{B, C}	2
{B, E}	3
{C, E}	2

C_2

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

2nd scan

C_2

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

C_3

Itemset	sup
{B, C, E}	2

3rd scan

L_3

Itemset	sup
{B, C, E}	2

Create initial 1-itemsets

Add each item to the initial list of candidate itemsets

```
def createC1(dataset):  
    "Create a list of candidate item sets of size one."  
    c1 = []  
    for transaction in dataset:  
        for item in transaction:  
            if not [item] in c1:  
                c1.append([item])  
    c1.sort()  
    #frozenset because it will be a ket of a dictionary.  
    return list(map(frozenset, c1))
```

Sort and return as list of sets

Identify itemsets that meet the support threshold

Calculate support counts
for each candidate

```
def scan(dataset, candidates, min_support):  
    "Returns all candidates that meets a minimum support level"  
    ssCnt = {}  
    for tid in dataset:  
        for can in candidates:  
            if can.issubset(tid):  
                ssCnt.setdefault(can, 0)  
                ssCnt[can] += 1  
  
    num_items = float(len(dataset))  
    retlist = []  
    support_data = {}  
    for key in ssCnt:  
        support = ssCnt[key] / num_items  
        if support >= min_support:  
            retlist.insert(0, key)  
            support_data[key] = support  
    return retlist, support_data
```

Check whether candidates
meet threshold

Generate the next list of candidates

(k-1)-itemsets

Iterate through all
pairs of itemsets

```
def aprioriGen(freq_sets, k):  
    "Generate the joint transactions from candidate sets"  
    retList = []  
    lenLk = len(freq_sets)  
    for i in range(lenLk):  
        for j in range(i + 1, lenLk):  
            L1 = list(freq_sets[i])[:k - 2]  
            L2 = list(freq_sets[j])[:k - 2]  
            L1.sort()  
            L2.sort()  
            if L1 < L2:  
                retList.append(freq_sets[i] | freq_sets[j]) # / is set union  
    return retList
```

Check whether pairs
differ by a single item

A|B returns the
union of A and B

Generate all Frequent Itemsets

Initialise L with frequent 1-itemsets

```
def apriori(dataset, min_support=0.5):  
    "Generate a list of candidate item sets"  
    C1 = createC1(dataset)  
    D = list(map(set, dataset))  
    L1, support_data = scanD(D, C1, min_support)  
    L = [L1]  
    k = 2  
    while (len(L[k - 2]) > 0):  
        Ck = aprioriGen(L[k - 2], k)  
        Lk, supK = scanD(D, Ck, min_support)  
        support_data.update(supK)  
        L.append(Lk)  
        k += 1  
  
    return L, support_data
```

While the list of (k-1)-itemsets is non-empty:

Generate candidate k-itemsets

Identify frequent k-itemsets

Keep frequent k-itemsets

Exercise: Generating Frequent Itemsets

- Frequent itemset generation
 -  code cell after “Generate frequent itemsets”
 -  code cell after “Itemset generation on sample data”
- Exploring support
 - Try different support thresholds
 - What’s a reasonable value for real data?
 - Do you have datasets that resemble transactions?
 - What about the apps/websites you use?

Rule Generation

- Given a frequent itemset L , find all non-empty subsets $x \subset L$ such that $x \rightarrow L - x$ satisfies the minimum confidence requirement
 - If $\{A,B,C,D\}$ is a frequent itemset, candidate rules:

$ABC \rightarrow D, ABD \rightarrow C, ACD \rightarrow B, BCD \rightarrow A,$

$A \rightarrow BCD, B \rightarrow ACD, C \rightarrow ABD, D \rightarrow ABC,$

$AB \rightarrow CD, AC \rightarrow BD, AD \rightarrow BC, BC \rightarrow AD,$

$BD \rightarrow AC, CD \rightarrow AB,$

- If $|L| = k$, then there are $2^k - 2$ candidate association rules


Efficient Rule generation

- How to efficiently generate rules from frequent itemsets?
- Similar to support, the confidence of rules generated from the same itemset has an anti-monotone property
- Example: $X = \{A,B,C,D\}$:
 - $c(ABC \rightarrow D) \geq c(AB \rightarrow CD) \geq c(A \rightarrow BCD)$
- Confidence is anti-monotone w.r.t. number of items on the RHS of the rule

Rule Generation for Apriori Algorithm


Low Confidence
Rule

Pruned
Rules



Association Rule

- From the example before, we had a frequent itemset: $\{B, C, E\}$
- From this frequent itemset $\{B, C, E\}$, we can have the following association rules



Exercise: Association Rule Mining

- Frequent itemset generation
 -  code cell after “Mine association rules”
 -  code cell after “Rule mining on sample data”
- Exploring confidence
 - Try different confidence thresholds
 - What’s a reasonable value for real data?
 - Can we use this for recommendation (e.g., Amazon, Netflix)?
- Using mlxtend library
 -  code cell after “Association analysis using mlxtend library”
 -  code cell after “Load the supermarket transaction datasets”
 - Mine association rules on Groceries datasets

Association Rule Mining with FP-Growth

Performance Bottlenecks of Apriori

- Bottlenecks of *Apriori*:
 - Candidate generation:
 - Generate huge candidate sets
 - Multiple scans of database

Overview of FP-Growth:

- Compress a large database into a compact, *Frequent-Pattern tree* (FP-tree) structure
 - Highly compacted, but complete for frequent pattern mining
 - Avoid costly repeated database scans
- Develop an efficient, FP-tree-based frequent pattern mining method (FP-growth)
 - A divide-and-conquer methodology: decompose mining tasks into smaller ones
 - Avoid candidate generation

Construct FP-tree

Two Steps:

1. Scan the transaction DB for the first time, find frequent items (single item patterns) and order them into a list **L** in frequency descending order.
2. For each transaction, order its frequent items according to the order in **L**; Scan DB the second time, construct FP-tree by putting each **frequency ordered transaction** onto it.

FP-tree Example: step 1

Step 1: Scan DB for the first time to generate **L**

Min_sup = 2

TID	Items bought
1	A, C, D
2	B, C, E
3	A, B, C, E
4	B, E



Item	Frequency
A	2
B	3
C	3
D	1
E	3


FP-tree Example: step 2

Step 2: scan the DB for the second time, order frequent items in each transaction

TID	Items bought	Ordered frequent items
1	A, C, D	C, A
2	B, C, E	C, B, E
3	A, B, C, E	C, B, E, A
4	B, E	B, E

FP-tree Example: step 2

Step 2: construct FP-tree



NOTE: Each transaction corresponds to one path in the FP-tree

Mining Frequent Patterns Using FP-tree

Starting the processing from the end of list **L**:

Step 1:


Construct **conditional pattern base** for each item in the header table

Step 2:

Construct **conditional FP-tree** from each conditional pattern base

Step 1: Construct Conditional Pattern Base

- Starting at the bottom of frequent-item header table in the FP-tree
- Traverse the FP-tree by following the link of each frequent item
- Accumulate all of **transformed prefix paths** of that item to form a **conditional pattern base**



Conditional pattern bases

Item	Cond. Pattern base
A	EBC:1, C:1
E	BC:2, B:1
B	C:2, $\{\}$
C	$\{\}$

Step 2: Construct Conditional FP-tree

- For each pattern base
 - Accumulate the count for each item in the base
 - Construct the **conditional FP-tree** for the frequent items of the pattern base

A-cond. pattern base:		A-conditional FP-tree	Frequent patterns relate to A:
EBC:1, C:1	→	{ (C:2) } A	AC:2

Step 2: Construct Conditional FP-tree

E-cond. pattern base: *E-conditional* FP-tree Frequent patterns relate to *E*:
BC:2, B:1 $\rightarrow \{ (B:3, C:2, BC:2) \} \mid E$ *EB :3, EC:2, BCE:2*

B-cond. pattern base: *B-conditional* FP-tree Frequent patterns relate to *B*:
C:2 $\rightarrow \{ (C:2) \} \mid B$ *BC:2*

All Frequent Item sets: { A:2, E:3, B:3, C:3,
BC:2, EC:2, AC:2, EB:3,
BCE:2 }

Exercise: Association Rule Mining Using FP-Growth

- Frequent itemset generation
 -  code cell after “Rules generation”
- Mine association rules using FP-growth on Groceries datasets
 - Try different confidence thresholds
 - What’s a reasonable value for real data?

Review

Today: Data Mining – Association Rule Mining

Objective

Learn techniques for unsupervised learning, with tools in Python.

Lecture

- Association rule mining

Readings

- Intro to Data Mining, Ch. 6
<http://www-users.cs.umn.edu/~kumar/dmbook/ch6.pdf>
- Intro to Data Mining, Ch. 8
<http://www-users.cs.umn.edu/~kumar/dmbook/ch8.pdf>
- Data Science from Scratch, Ch. 11 & 19

Exercises

- Associations from scratch
- Associations using mlxtend
- Associations using pyfpgrowth

TODO in W5

- define experimental framework for A2

Additional Reading (not examinable)

- Tan et al. Introduction to data mining.
<https://goo.gl/hWwuZb>
- Aggarwal. Data mining: the textbook.
<https://goo.gl/IQqLwT>
- Han. Data mining: concepts and techniques.
<https://goo.gl/CFIMMs>
- Scikit-learn user guide, § 2 (Unsupervised learning).
http://scikit-learn.org/stable/unsupervised_learning.html

Other tools and Techniques (not examinable)

- Scikit-learn user guide, § 4.4 (Dimensionality reduction).
http://scikit-learn.org/stable/modules/unsupervised_reduction.html
- Scikit-learn user guide, § 2.7 (Outlier detection).
http://scikit-learn.org/stable/modules/outlier_detection.html
- Etc

Project and discussion time

*Time for you to talk
to tutors, instructors and each other.*